

IBADAT ALI

ibadcodes@gmail.com | +92 322 069 2321 | linkedin.com/in/mirzaibadatali | github.com/Ibadat-Ali86

ibadat-ali-portfolio.vercel.app | huggingface.co/ibadatali | fiverr.com/ibadcodes

Lahore, Pakistan | Available for Onsite & Remote Roles

PROFESSIONAL SUMMARY

AI/ML Engineer with industry experience across internship, full-time employment, and independent freelance delivery — building production-deployed deep learning systems, LLM-integrated automation pipelines, and end-to-end client solutions with measurable business outcomes. Shipped a 4-model forecasting ensemble (MAPE 1.23%, 580K+ transactions), industrial predictive maintenance automation (\$38K attributed cost savings, RMSE 14.37), and multimodal clinical AI workflows — all with explicit fallback layers, confidence gates, and audit-grade traceability producing zero silent failures. Fluent across the full MLOps lifecycle from model training through FastAPI deployment, Docker packaging, and production monitoring; delivers zero-dependency client handoffs backed by structured documentation.

TECHNICAL SKILLS

Machine Learning & Deep Learning: PyTorch, Scikit-Learn, XGBoost, LightGBM, Prophet, SARIMA, LSTM, TCN, Isolation Forest, Autoencoder, TensorFlow, Knowledge Distillation

LLM & Agents: Claude API, Gemini API, OpenAI API, LangChain, MCP (Model Context Protocol), RAG Pipeline Design, Prompt Engineering, Structured Output, Semantic Search

Data Engineering: Python, Pandas, NumPy, SQL, PostgreSQL, MySQL, SQLite, SQLAlchemy, Prisma, ETL Pipeline Design

MLOps & Deployment: FastAPI, Flask, Docker, MLflow, GitHub Actions, Vercel, Hugging Face Spaces, n8n, Electron, REST APIs

Explainability & Evaluation: SHAP Attribution, Confidence Gates, Audit Logging, Human-in-the-Loop Review, Model-Boundary Annotations, A/B Testing

Visualization & Reporting: Matplotlib, Seaborn, Plotly, Power BI, Automated PDF/CSV Report Generation

EXPERIENCE

ML/AI Engineer | Full-Time | On-site

Dec 2025 – Mar 2026

Soft Shack

- Designed and deployed production ML/AI pipelines for client-facing automation systems, integrating LLM APIs (OpenAI, Gemini) with FastAPI serving layers and Docker containerization for reproducible deployment.
- Built end-to-end AI agent workflows and RAG-based retrieval systems; delivered production systems with audit logging and confidence gates to ensure zero silent failures across inference pipelines.
- Collaborated cross-functionally to translate client requirements into model architectures; produced structured handoff documentation enabling client engineers to operate systems independently post-delivery.

Data Scientist Intern | Internship | 4 Months

May 2025 – Aug 2025

Arch Technologies

- Developed and evaluated machine learning models for business forecasting and predictive analytics use cases using Python, Scikit-Learn, XGBoost, and Pandas; delivered model performance reports with SHAP-based feature attribution.
- Engineered ETL pipelines for data ingestion and preprocessing from SQL and CSV sources; applied statistical inference and data validation to ensure model-ready data quality; produced stakeholder-facing analysis outputs and visualizations.

Independent AI/ML Freelancer | Freelance | Remote

2025 – Present

Fiverr (Self-Employed)

- Delivered production AI/ML solutions to international clients across two service tracks: (1) ML/Deep Learning/Data Science systems (predictive modeling, forecasting, NLP, RAG, LLM integration, computer vision with YOLOv8/OpenCV, FastAPI deployment, Docker); (2) AI automation workflows (n8n, LangChain, AI agents, OpenAI API, WhatsApp Business API integration).
- Delivered systems with explicit fallback logic, structured handoff documentation, and zero ongoing developer dependency; production benchmarks include MAPE 1.23% forecasting accuracy and \$38K attributed cost savings.

SELECTED PROJECTS

SentinelIQ — Industrial Predictive Maintenance Automation Live | Open Source

[Source](#)

Stack: PyTorch, TCN, LSTM, Isolation Forest, Autoencoder, SHAP, PuLP, FastAPI, Docker

- Engineered TCN-LSTM architecture in PyTorch for equipment remaining-useful-life prediction (RMSE 14.37), integrated with PuLP optimizer to translate predictions into optimal maintenance windows — attributing \$38K in cost savings with SHAP per-decision attribution for full operational auditability.
- Deployed dual anomaly detection layer (Isolation Forest + Autoencoder) as independent validation; conflicting signals trigger human review flags rather than auto-scheduling, eliminating silent failure modes by architecture.

AdaptIQ — Automated Business Demand Forecasting SaaS Live

Stack: Prophet, XGBoost, SARIMA, LSTM, FastAPI, Docker, Hugging Face Spaces

- Built 4-model ensemble (Prophet, XGBoost, SARIMA, LSTM) over 580K+ retail transactions achieving MAPE 1.23% with automated regime-aware model selection and fallback logic — zero manual analyst override required.
- Architected confidence-interval scenario planning and automated PDF/CSV report generation via FastAPI; non-technical clients consume production forecasts without any data science dependency.

CareVision — Multimodal Clinical AI Workflow Google Gemma 4 Hackathon 2025 | Live

[Source](#)

Stack: Gemini 2.0 Flash, FastAPI, PostgreSQL, SQLAlchemy, React 18, TypeScript

- Integrated Gemini 2.0 Flash into a clinical PWA with model-boundary annotations and confidence gates; step-level audit logging produces zero silent failures across the entire inference workflow.
- Implemented structured output schema enforcement and offline-capable inference fallback backed by full PostgreSQL workflow traceability.

CodeScope MCP Preflight — Agent Tool-Calling & RAG Retrieval OpenAI Build Week 2026 | Devpost Source | Devpost

Stack: Python, MCP, sentence-transformers, Chroma, Pydantic, pytest, uv, Tree-sitter

- Built local-first, read-only MCP server exposing four structured tools to AI coding agents — repository inventory, behavioral code search, exact symbol lookup, and similar-implementation detection — enabling agents to make an explicit REUSE, EXTEND, or CREATE decision before editing.
- Validated through controlled comparison testing: CodeScope surfaced authoritative production symbols and unsafe deprecated alternatives before code generation, making evidence behind patches explicit, source-grounded, and repeatable across both test runs without modifying source files.

n8n AI Automation Systems

Stack: n8n, Gemini 2.0 Flash, WhatsApp Business API, Google Sheets, Gmail

- Engineered lead generation pipeline (form trigger, Maps scraping, Gemini enrichment, Gmail outreach, explicit fallback routing — zero manual steps) and WhatsApp Restaurant Chatbot with Gemini AI Agent, session memory, and three-tool interface (FAQ, orders, inventory) — 2.61s pipeline, 1.13s agent inference, 655 tokens per turn.

Client Deliveries — Production Systems

- **Mr. Stationer (Offline-First POS):** Designed, built, and shipped production POS in React, Electron, Node.js, and SQLite for real retail client; transaction processing, inventory management, and tiered automatic backup operational without internet — client operates independently post-handoff.
- **Evershaheen Academy (Enterprise LMS):** Delivered multi-campus school management system ([evershaheenacademy.com](#)) in Next.js 15, Prisma, and PostgreSQL; nine-layer technical documentation enables AI coding agents and client engineers to extend the system without developer involvement.

RESEARCH

Co-Author — Manuscript in Preparation (2025–2026)

TopoLite-KD: Efficient Topology-Aware Knowledge Distillation for COVID-19 CT-Slice Classification

GIFT University, Gujranwala, Pakistan | Supervised by: [CO-AUTHOR / ADVISOR NAME] | [Repository](#)

EDUCATION

Bachelor of Science in Data Science

May 2027

GIFT University, Gujranwala, Pakistan

Relevant Coursework: Machine Learning, Deep Learning, Database Systems, Statistical Inference, Data Structures & Algorithms, Operations Research

Research: Co-author, *TopoLite-KD: Efficient Topology-Aware Knowledge Distillation for COVID-19 CT-Slice Classification* (manuscript in preparation)

CERTIFICATIONS

Claude 101 Certification

Anthropic, 2026

AI Fluency Framework & Foundations

Anthropic, 2026

HACKATHONS & COMPETITIONS

- **OpenAI Build Week — Devpost (Jul 2026)** — Submitted [CodeScope MCP Preflight](#): local-first MCP server enabling AI coding agents to perform structured repository preflight (REUSE, EXTEND, or CREATE) before code generation; built with sentence-transformers, Chroma vector retrieval, and a read-only discovery boundary.
- **Google Gemma 4 Hackathon (2025)** — Submitted CareVision: multimodal clinical AI PWA with Gemini 2.0 Flash integration, explainability-first workflow design, confidence gates, and offline-capable inference architecture.

- **Digital Pakistan Speed Programming Competition — DPSPC (2025)** — Ranked in the Top 50 teams nationwide out of approximately 450 competing teams in a national speed programming and algorithmic problem-solving competition; challenges evaluated data structures, optimal time/space complexity, and problem-solving under strict time limits.
- **InnovateX 2.0 Coding Competition — GIFT University (2025)** — Competed in a senior-level algorithmic coding competition organized by Brackets GIFT University; solved competitive programming challenges emphasizing efficient algorithm design and optimal time/space complexity; recognized among top performers by industry evaluators.
- **AI/ML Content Creator** — Educational ML content on TikTok (@deepfx_ai) and Instagram (@explainable_ai); model architecture and applied AI concepts for technical and non-technical audiences.